



Introduction To **Azure Machine Learning**



Expected Learning Outcomes

Azure Machine Learning

- **Supervised Machine Learning**
- **Machine Learning Algorithms in Azure ML**
- **Workflow for Supervised Machine Learning**
- **Algorithm Selection**





Supervised

Machine Learning





➡ Unlabeled Data

➡ Labeled Data




Square Footage	Zip Code	Price
2000	98052	5000 \$
3000	75284	4000 \$
3500	68164	5500 \$
2000	45465	5400 \$
7520	98645	10000 \$





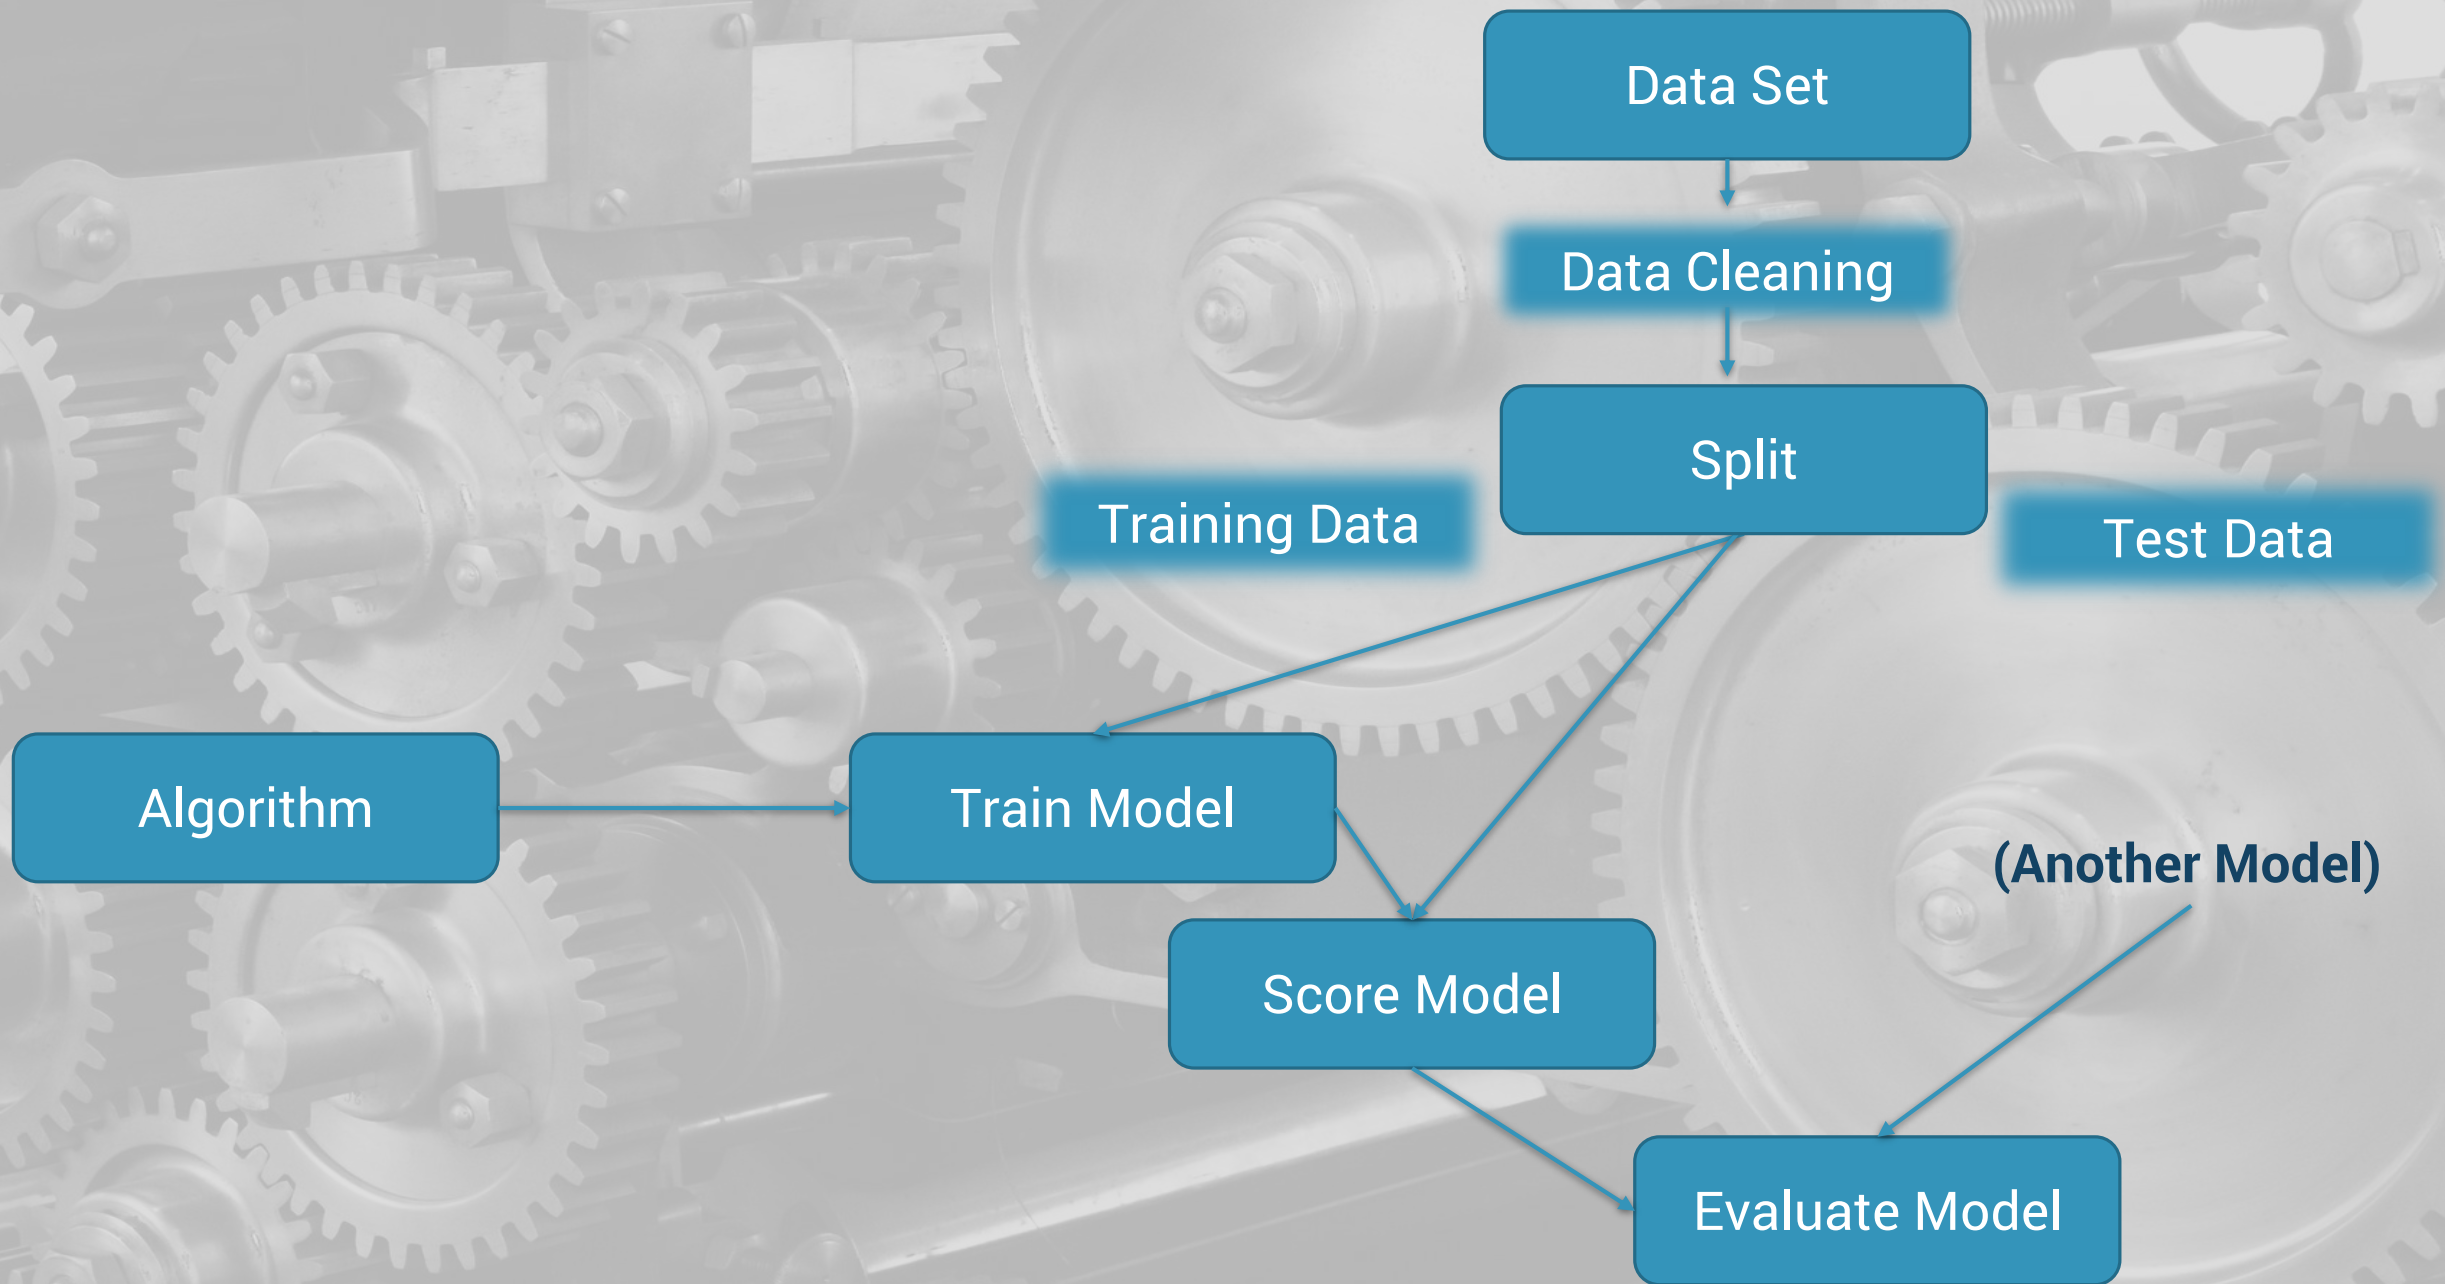
- **Classification** - predicting a discrete category
- **Regression** - predicting a value
- **Anomaly Detection** - identifying unusual data
- **Clustering** - grouping similar data together





- ➔ 25 different machine learning algorithms
- ➔ Modules for importing and cleaning data
- ➔ Create experiments to train and test models
- ➔ Publish models as web service (REST)







- ↪ Sample data
- ↪ Upload dataset from local machine
- ↪ Enter data manually
- ↪ Import data from external website





↪ Modules

↪ Datasets

↪ Experiments





Microsoft Azure Machine Learning: Algorithm Cheat Sheet

This cheat sheet helps you choose the best Azure Machine Learning Studio algorithm for your predictive analytics solution. Your decision is driven by both the nature of your data and the question you're trying to answer.

ANOMALY DETECTION

One-class SVM

>100 features,
aggressive boundary

PCA-based anomaly detection

Fast training

CLUSTERING

K-means

MULTI-CLASS CLASSIFICATION

Fast training, linear model

Multiclass logistic regression

Accuracy, long training times

Multiclass neural network

Accuracy, fast training

Multiclass decision forest

Accuracy, small memory footprint

Multiclass decision jungle

Depends on the two-class classifier, see notes below

One-v-all multiclass

REGRESSION

Ordinal regression

Data in rank ordered categories

Poisson regression

Predicting event counts

Fast forest quantile regression

Predicting a distribution

Linear regression

Fast training, linear model

Bayesian linear regression

Linear model, small data sets

Neural network regression

Accuracy, long training time

Decision forest regression

Accuracy, fast training

Boosted decision tree regression

Accuracy, fast training,
large memory footprint

Discovering structure
Finding unusual data points

START

Predicting values

Three or more
Predicting categories

Two

TWO-CLASS CLASSIFICATION

Two-class SVM

>100 features,
linear model

Two-class averaged perceptron

Fast training,
linear model

Two-class logistic regression

Fast training,
linear model

Two-class Bayes point machine

Fast training,
linear model

Accuracy,
fast training

Two-class decision forest

Accuracy,
fast training,
large memory footprint

Two-class boosted decision tree

Accuracy,
small memory footprint

Two-class decision jungle

>100 features

Two-class locally deep SVM

Accuracy, long training times

Two-class neural network

